


Many-to-one model

User-level threads (a.k.a green threads)



One kernel thread per process
thread management and scheduling is delegated to a library

- e.g pthreads `PTHREAD_SCOPE_PROCESS`
- e.g Java threads


➡ The kernel is not involved

- ✓ Very lightweight and fast
- ✓ All threads can be blocked if one of them is waiting on an event
- ✓ Cannot be scheduled on multiple cores

Many-to-many model

Hybrid threads

(a.k.a n:m threading)



User threads implemented on kernel threads

- e.g (old) Solaris

➡ Multiple kernel-level threads per process